

HBase Fundamentals

CISC-525

Phil Grim

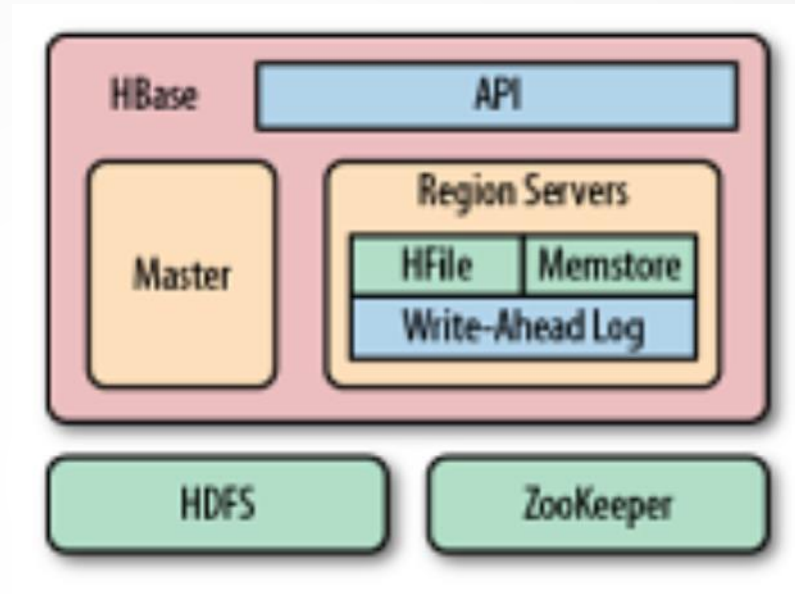
Overview

- ▶ HBase Architecture
- ▶ Data Storage

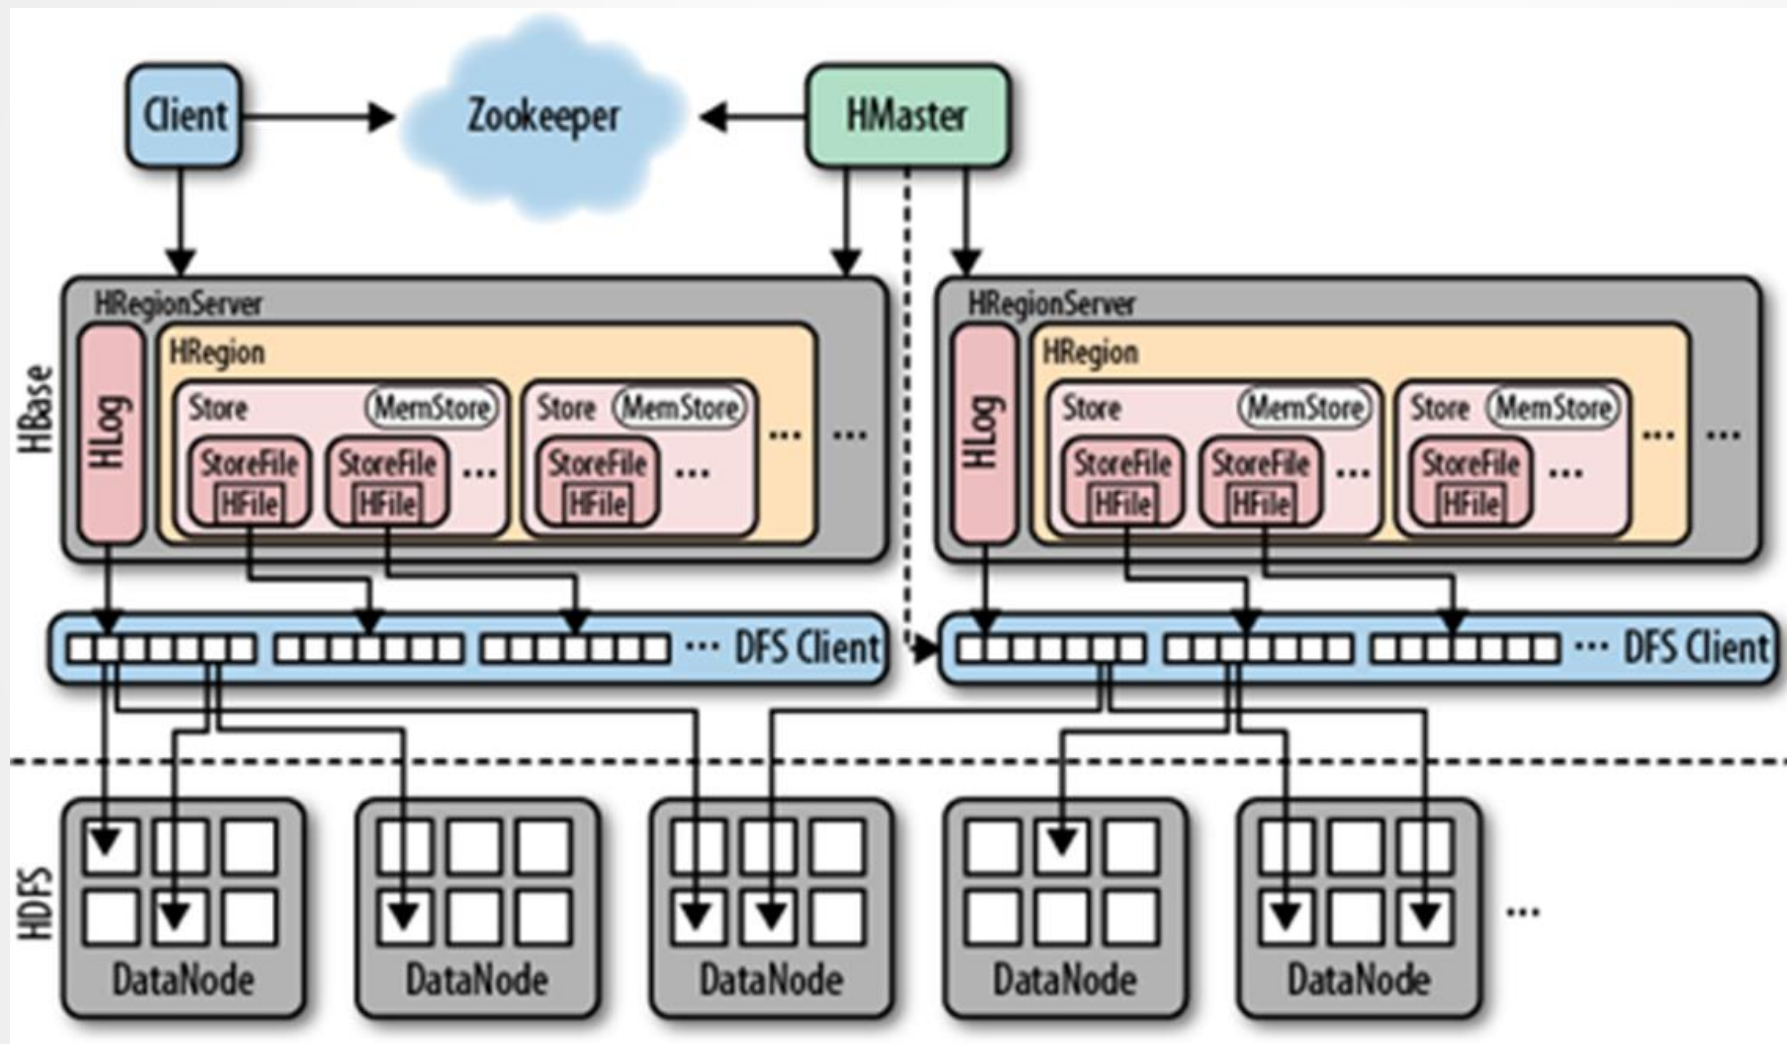


HBase Architecture

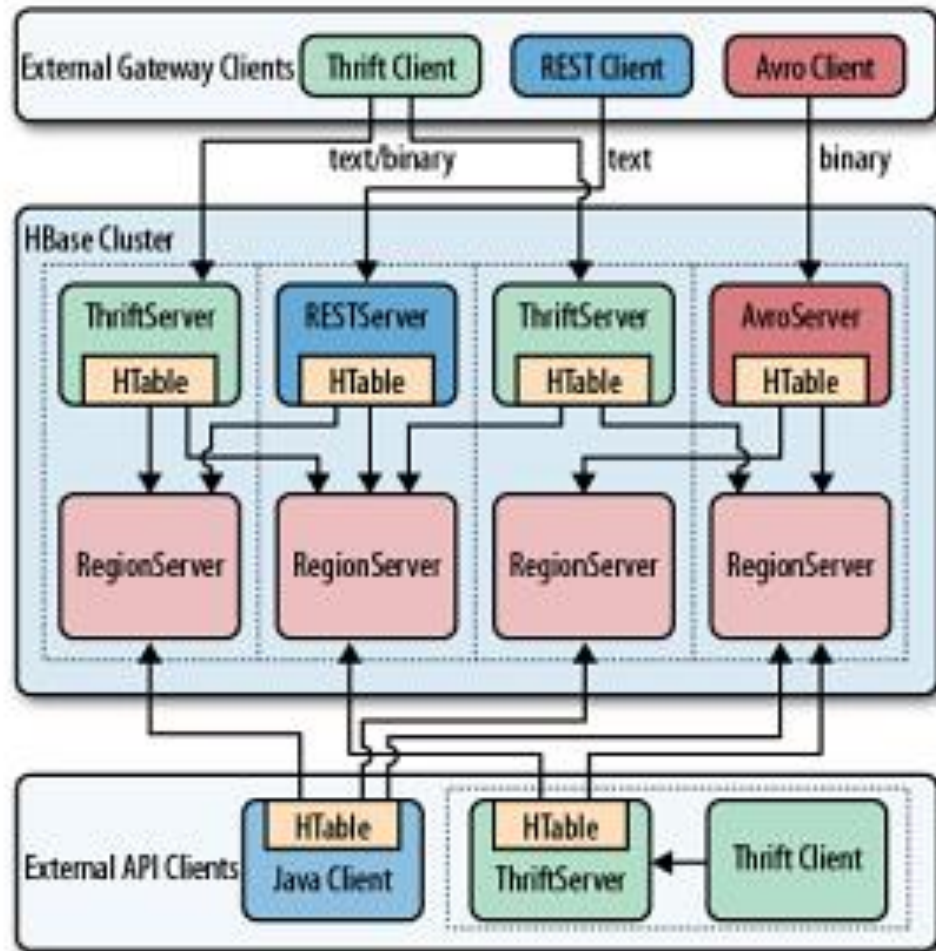
- ▶ Master/Slave
 - ▶ HMaster
 - ▶ RegionServer
- ▶ Coordinated with Zookeeper
- ▶ Data stored in HDFS



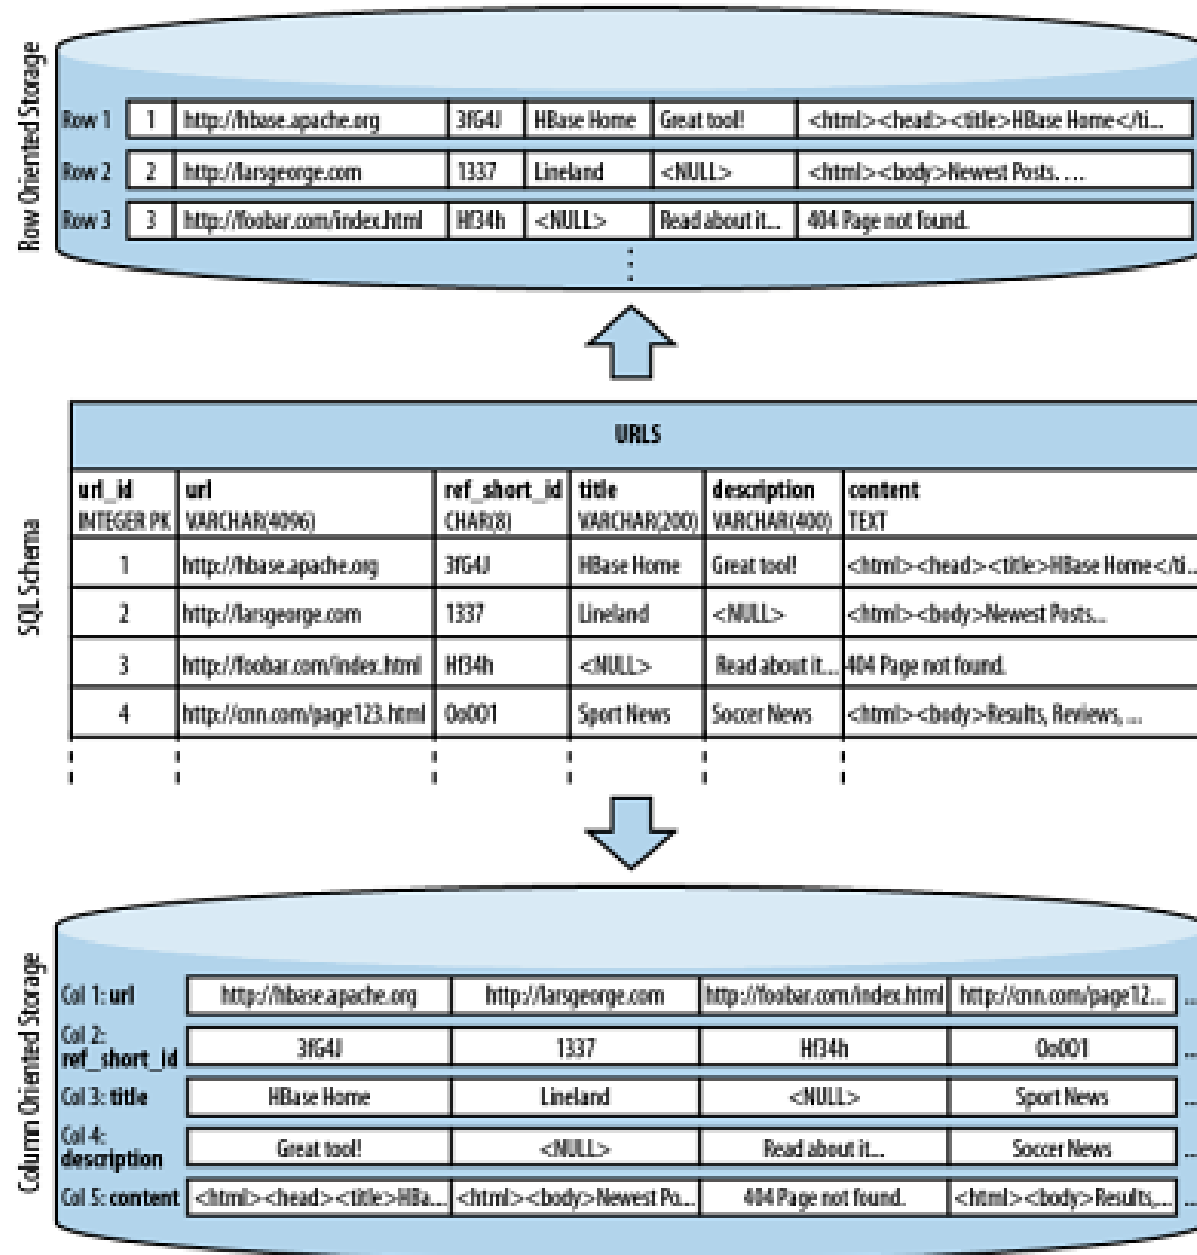
HBase Architecture



HBase Architecture

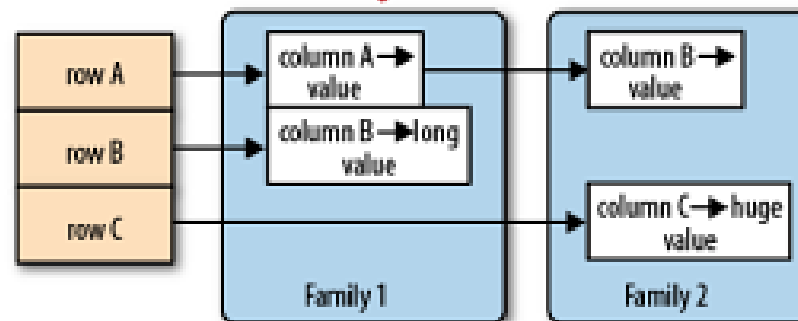


Data Storage

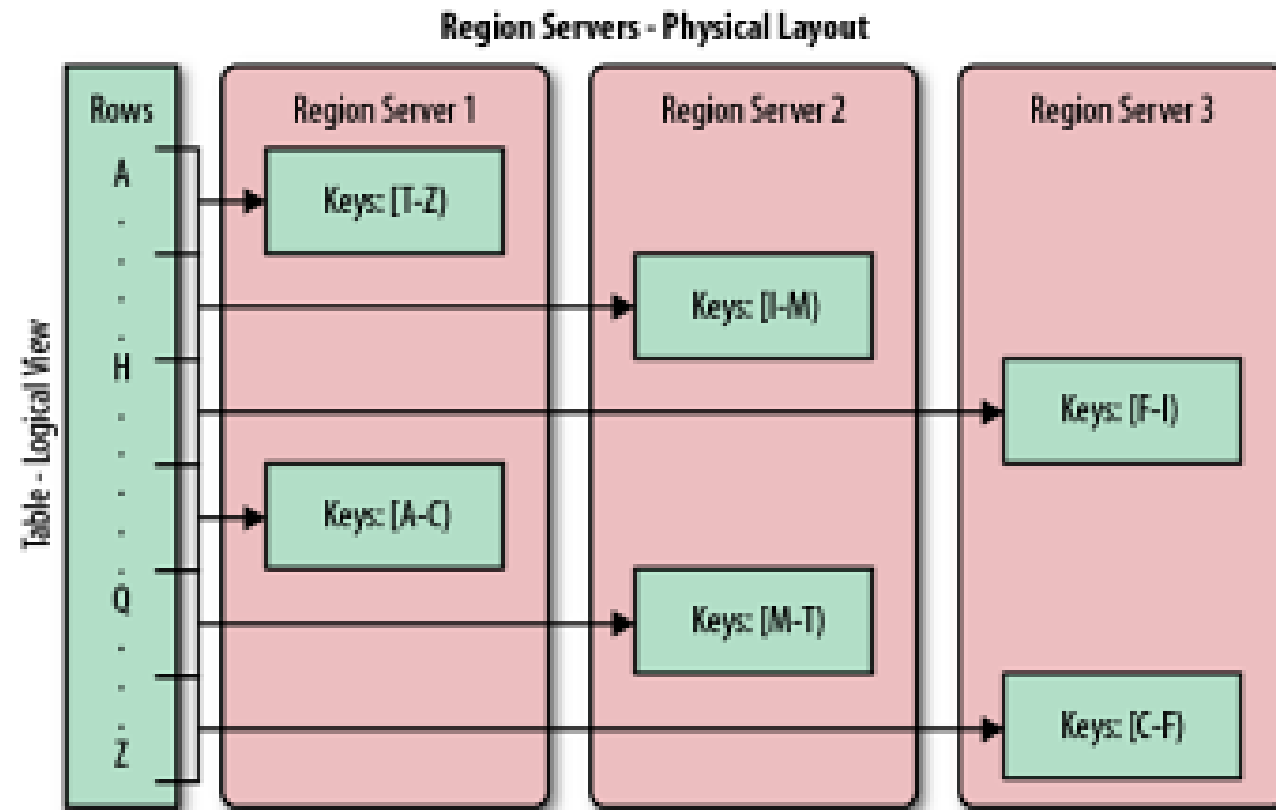


Data Storage

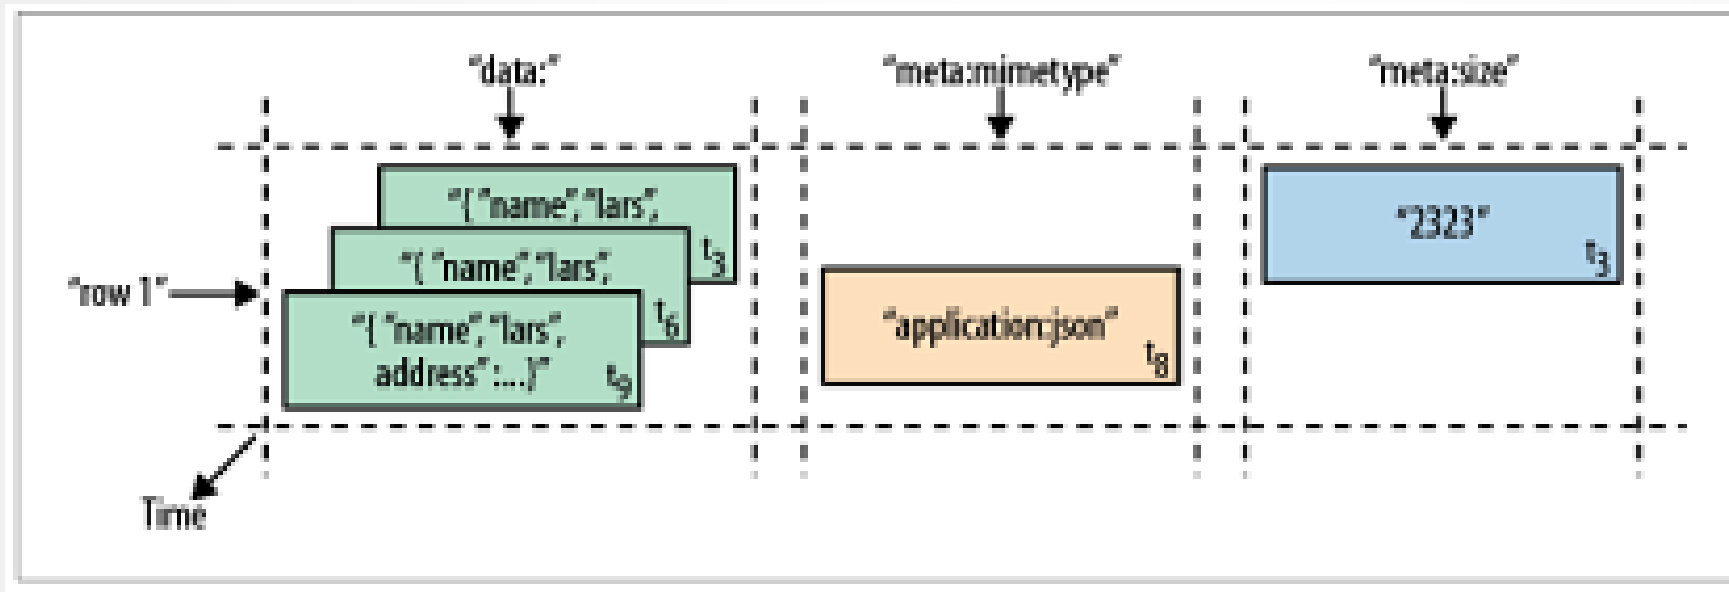
	column A (int)	column B (varchar)	column C (boolean)	column D (date)
row A				
row B				
row C			NULL?	
row D				



Data Storage



Data Storage



Continued Reading

HBase: The Definitive Guide

<http://shop.oreilly.com/product/0636920014348.do>

Apache HBase website

<http://hbase.apache.org/>

